

Data Science

Lecture 10-1: Image Data Processing
(Image Filtering)



UNIVERSITY
OF AMSTERDAM

Lecturer: Yen-Chia Hsu

Date: Mar 2023

This lecture covers basic concepts of image filtering and feature extraction.

There are many different types of tasks in [Computer Vision](#). This lecture will only cover the image/video classification, which gives only one label to an image/video.

Classification



CAT

No spatial extent

Semantic Segmentation



GRASS, CAT,
TREE, SKY

No objects, just pixels

Object Detection



DOG, DOG, CAT

Multiple Object

Instance Segmentation



DOG, DOG, CAT

[This image is CC0 public domain](#)

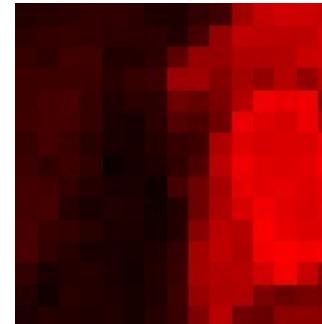
People can see images directly, but computers read only numbers. Typically, computers store images as **pixels** with RGB channels with values ranging from 0 to 255.



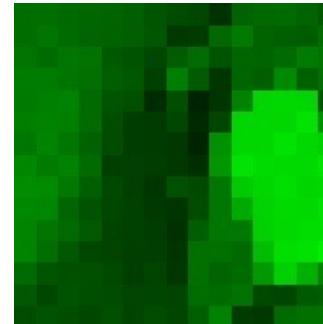
[Source](#)



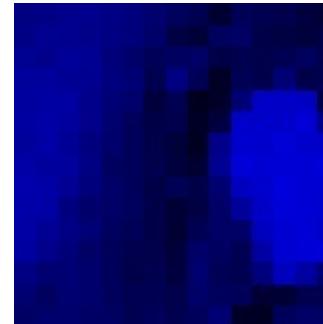
Image patch



Red



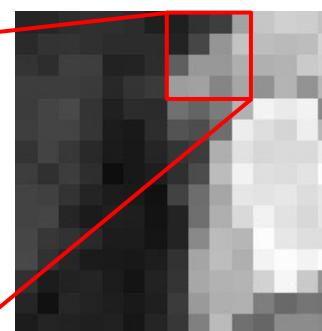
Green



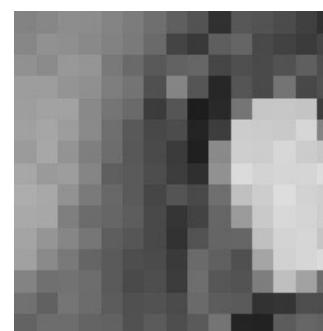
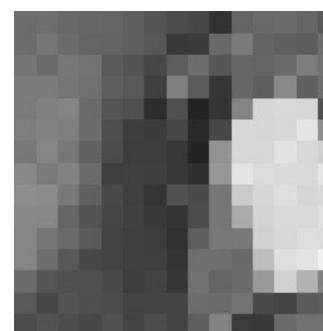
Blue

Colored visualization of the RGB channels

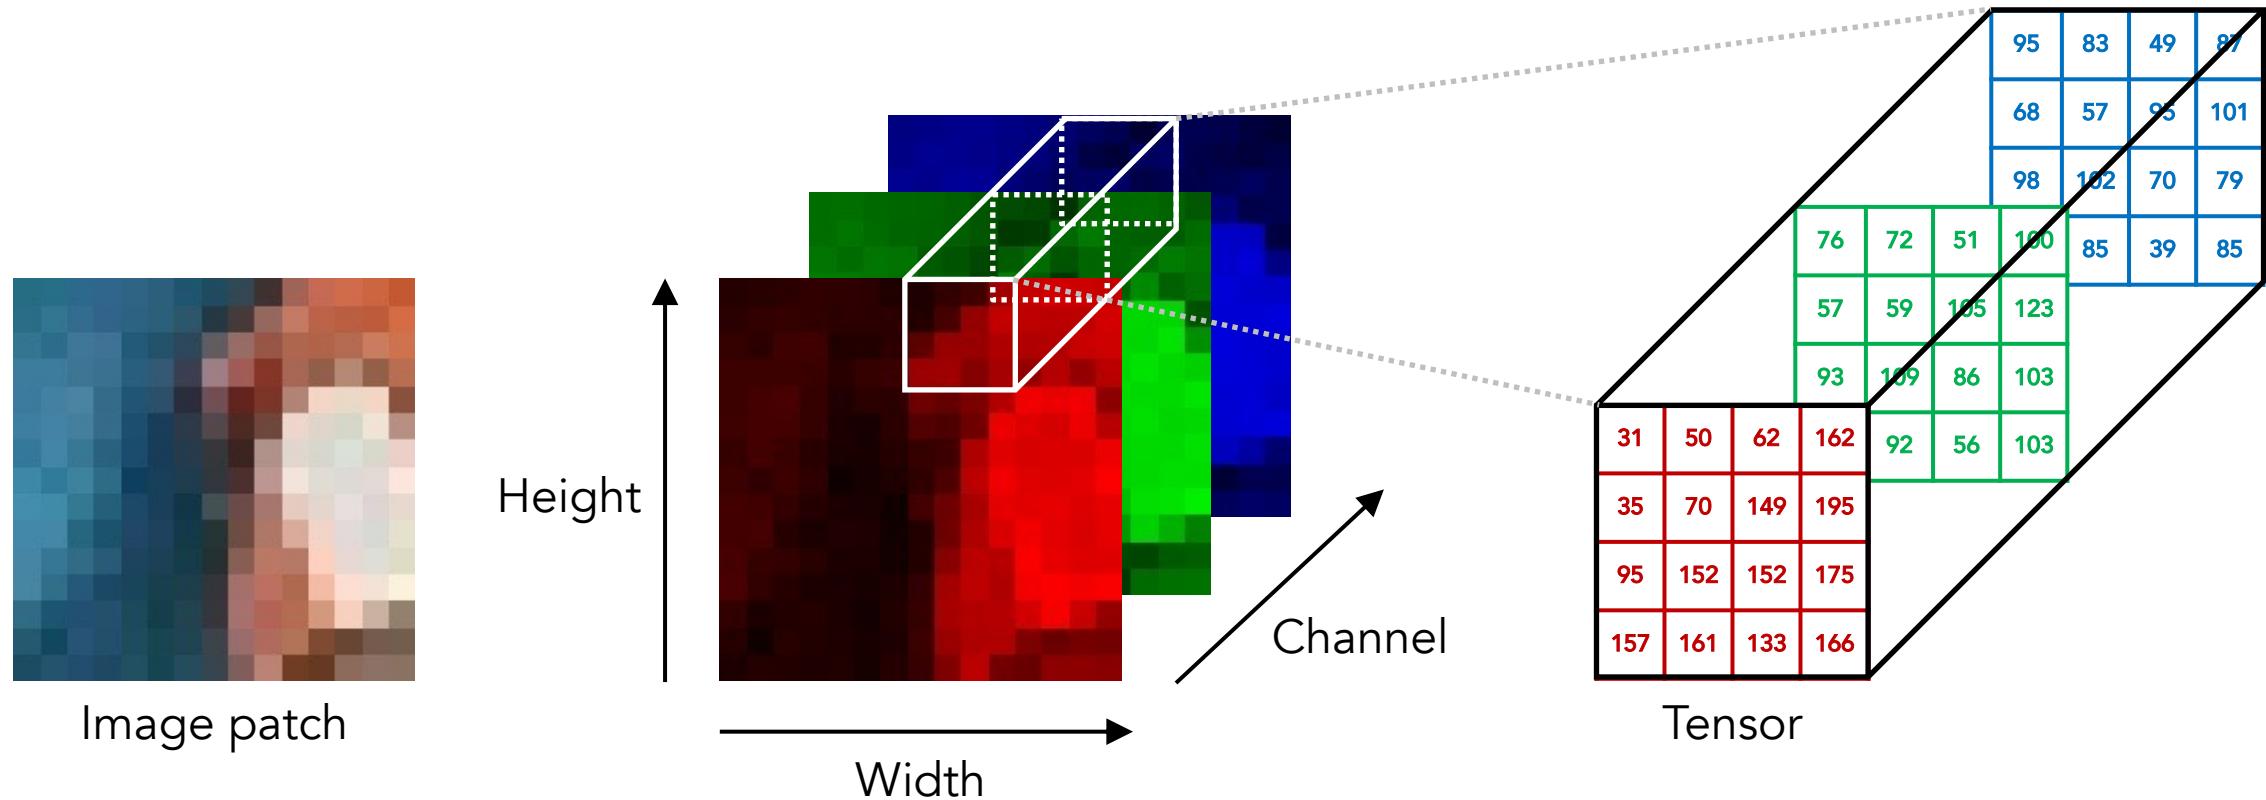
31	50	62	162
35	70	149	195
95	152	152	175
157	161	133	166



The intensity values (pixel values) per channel



Images can be represented as a 3D tensor (width*height*channels). For RGB images, we have 3 channels corresponding to the pixel intensity of red, green, and blue colors.



Before the deep learning era, Computer Vision used hand-crafted features. Researchers developed different image filters/kernels to extract features using **convolution**.



Input image

Discrete convolution: sum of the element-wise multiplication

Input				
7	2	3	3	8
4	5	3	8	4
3	3	2	8	4
2	8	7	2	7
5	4	4	5	4

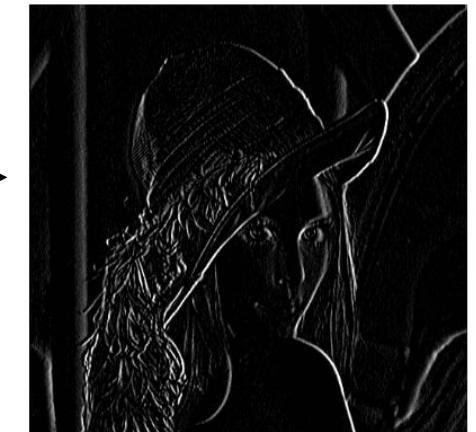
Filter
(Kernel)

1	0	-1
1	0	-1
1	0	-1

*

6		

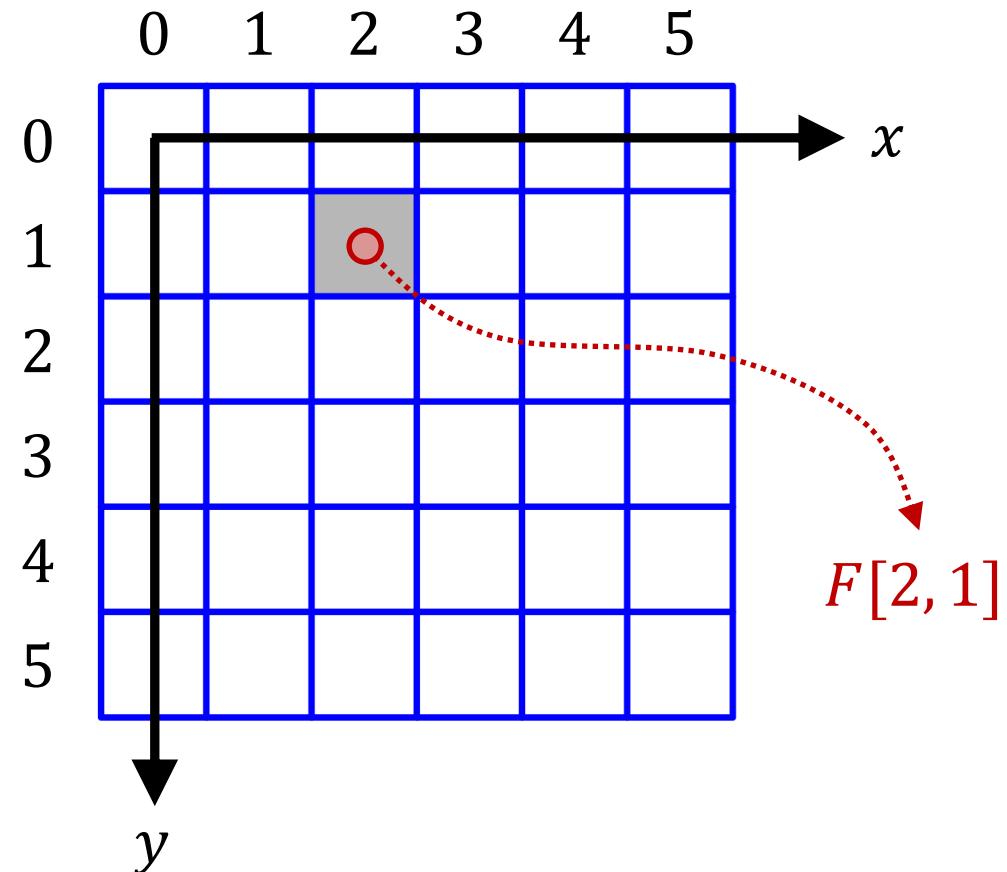
Output



Output image

$$\begin{aligned} & 7 \times 1 + 4 \times 1 + 3 \times 1 + \\ & 2 \times 0 + 5 \times 0 + 3 \times 0 + \\ & 3 \times -1 + 3 \times -1 + 2 \times -1 \\ & = 6 \end{aligned}$$

About image coordinates, the origin is at the top-left pixel. Notation $F[x, y]$ means the value of the center of the pixel for image F at location $[x, y]$ in the 2D array.

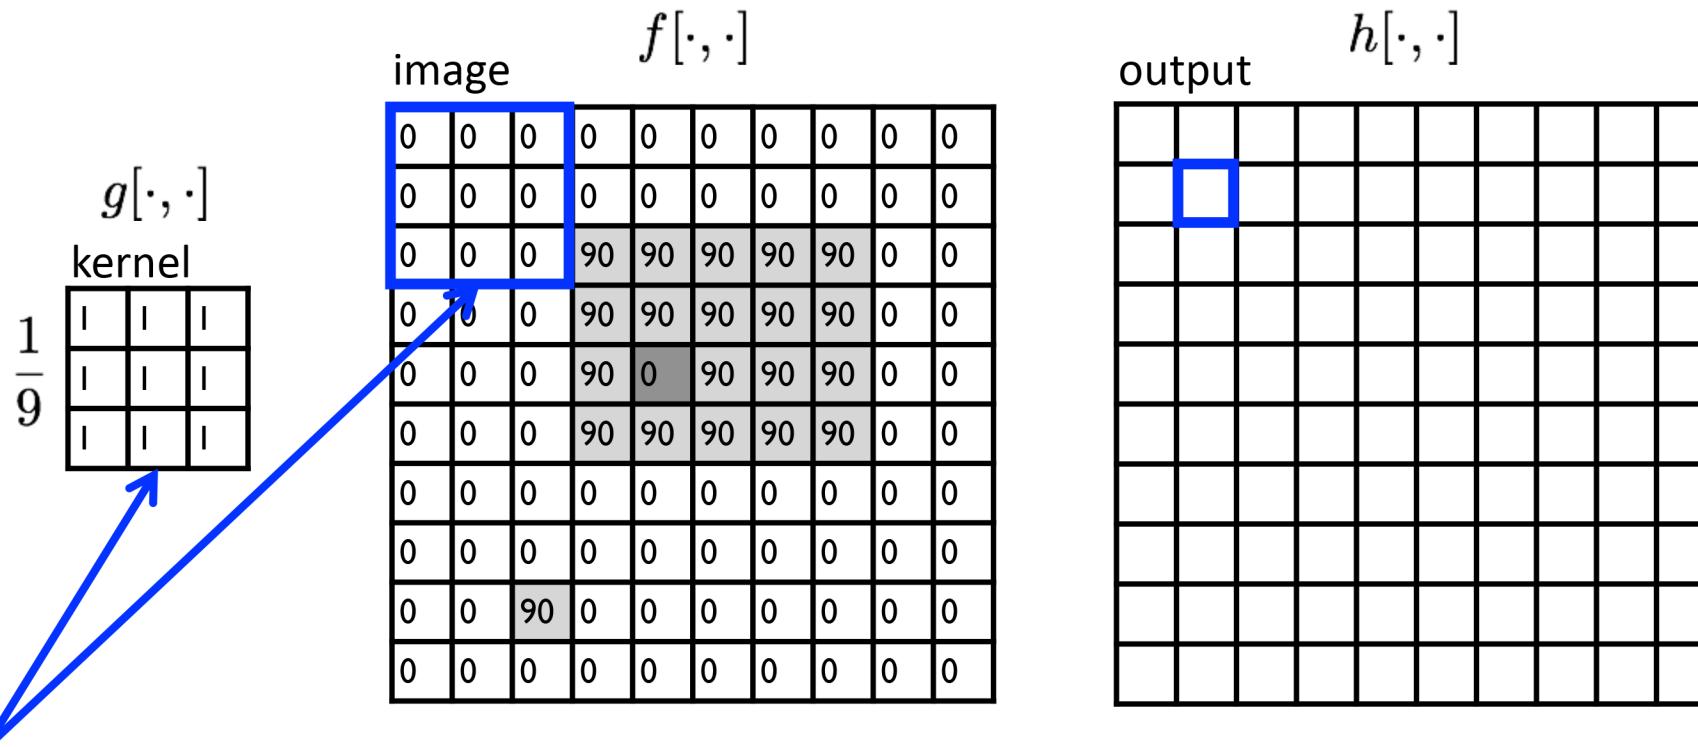


$F[x, y]$

F is an array of numbers (pixels)

x, y are integers (column/row indices)

Let's run the box filter

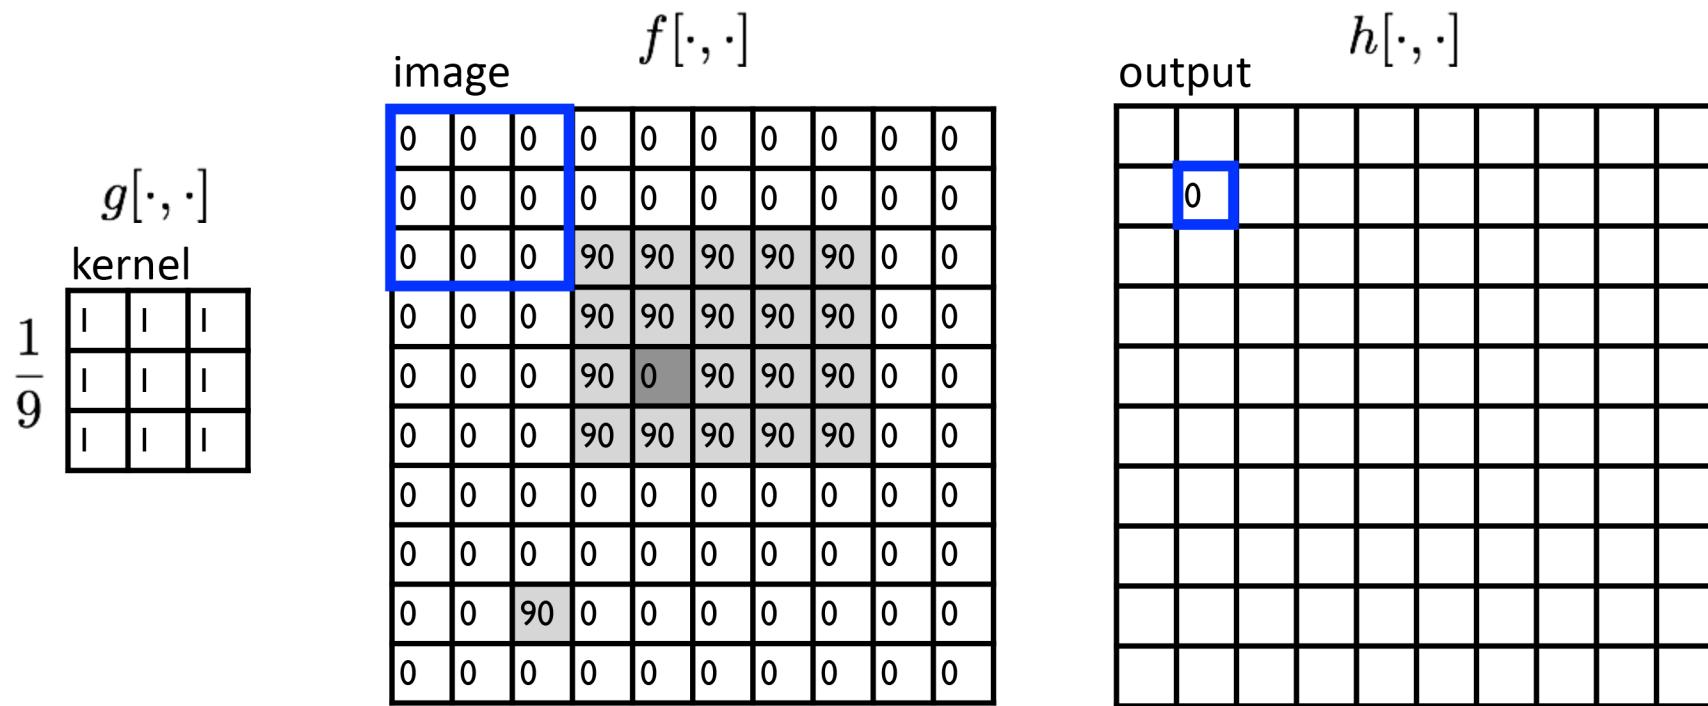


note that we assume that
the kernel coordinates
are centered

$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output filter image (signal)

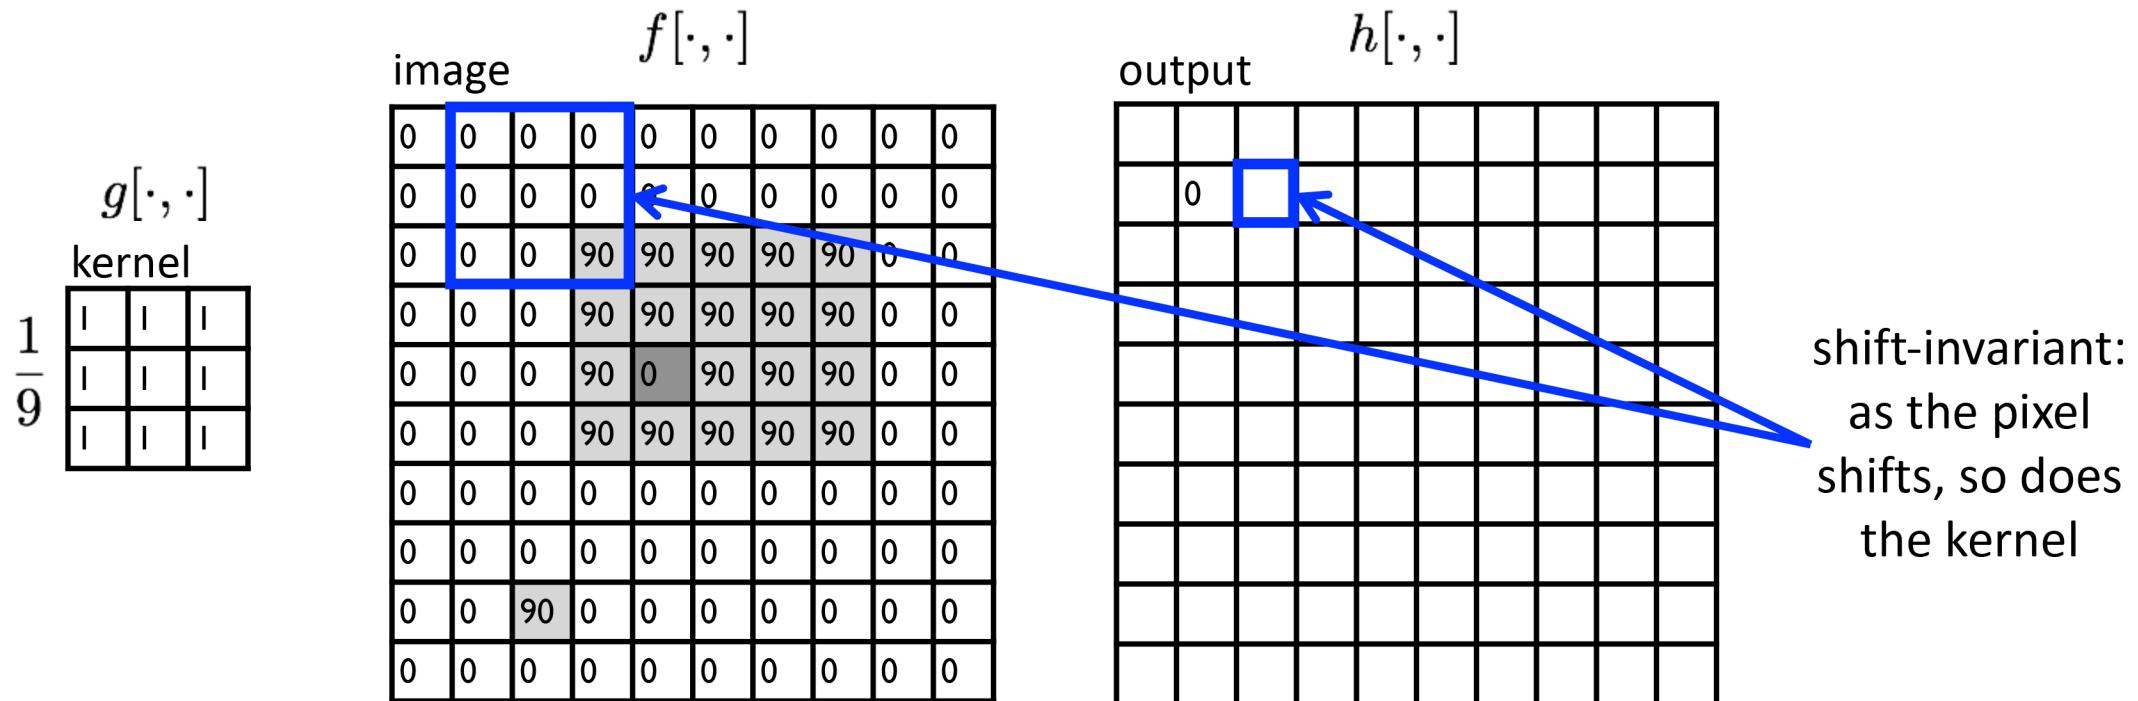
Let's run the box filter



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output filter image (signal)

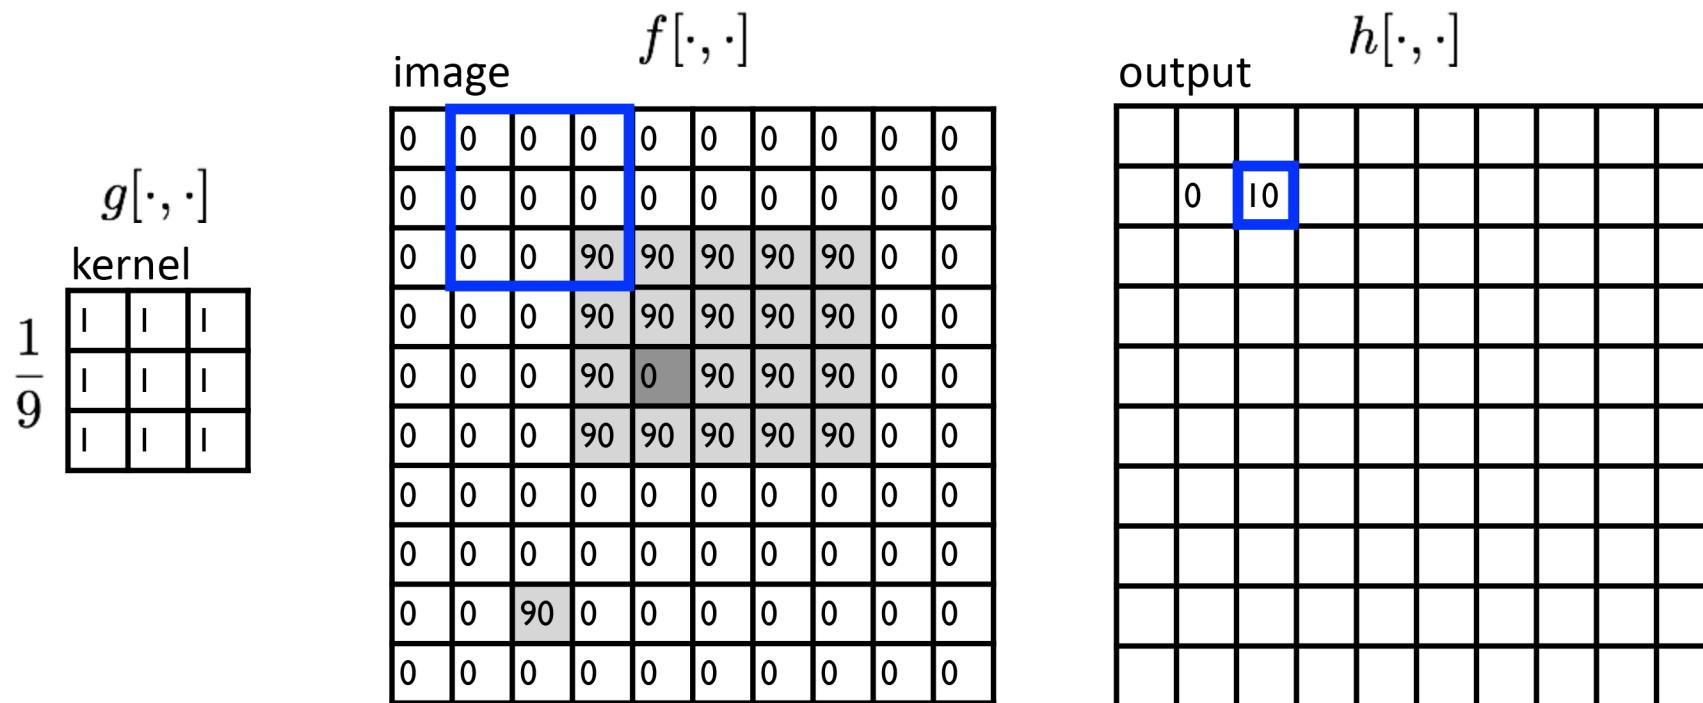
Let's run the box filter



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output filter image (signal)

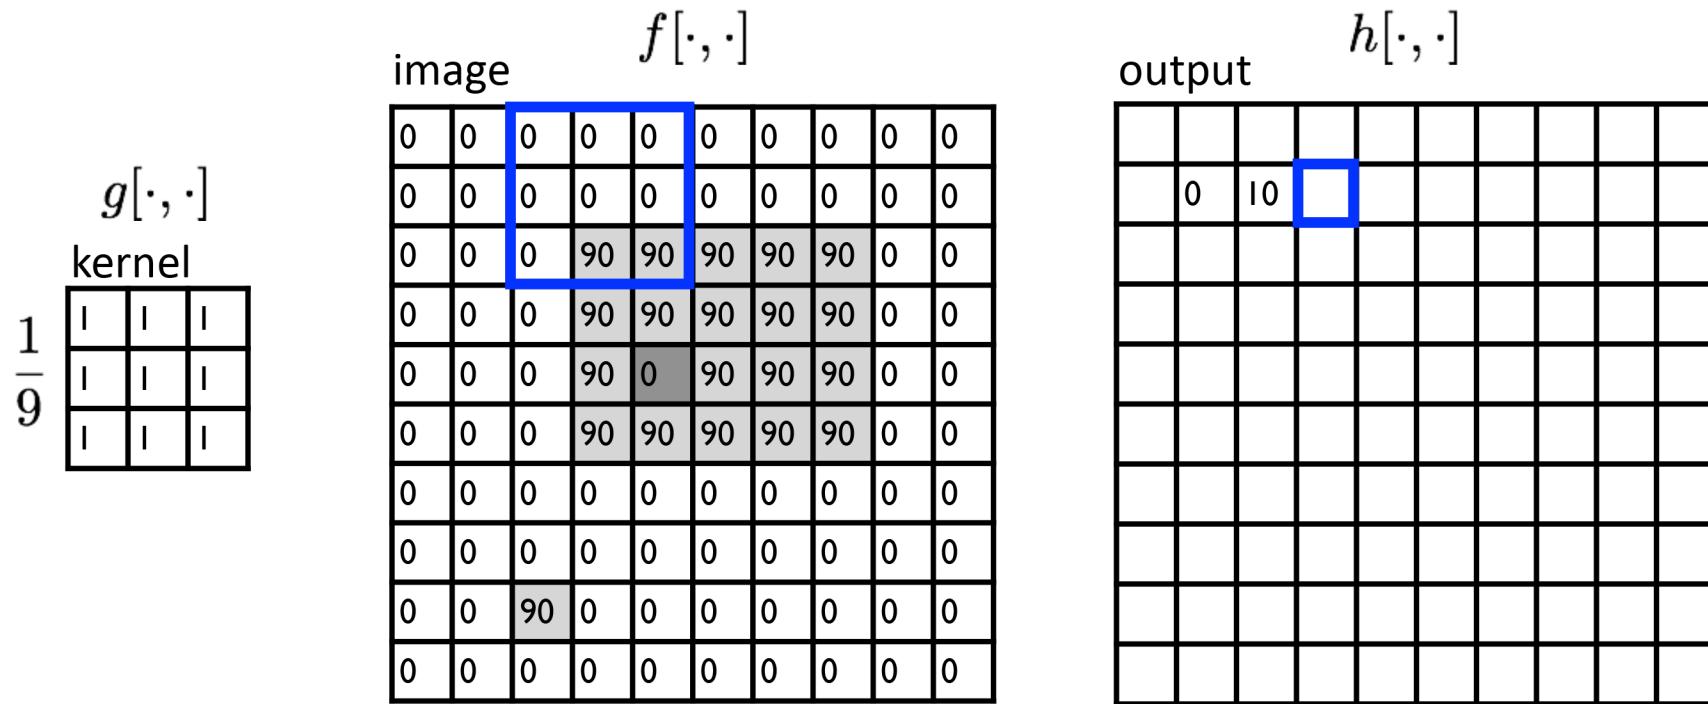
Let's run the box filter



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output filter image (signal)

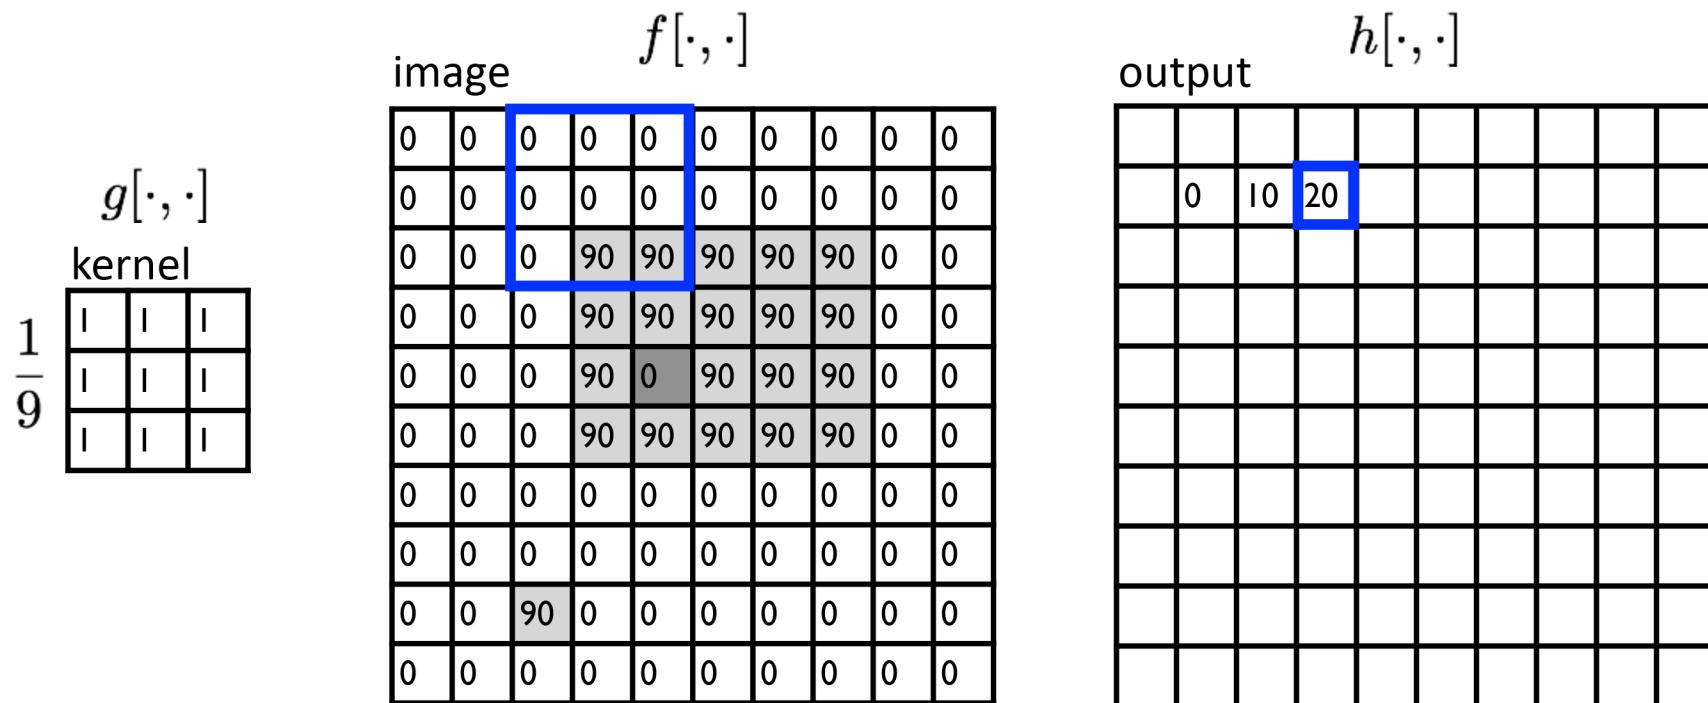
Let's run the box filter



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output filter image (signal)

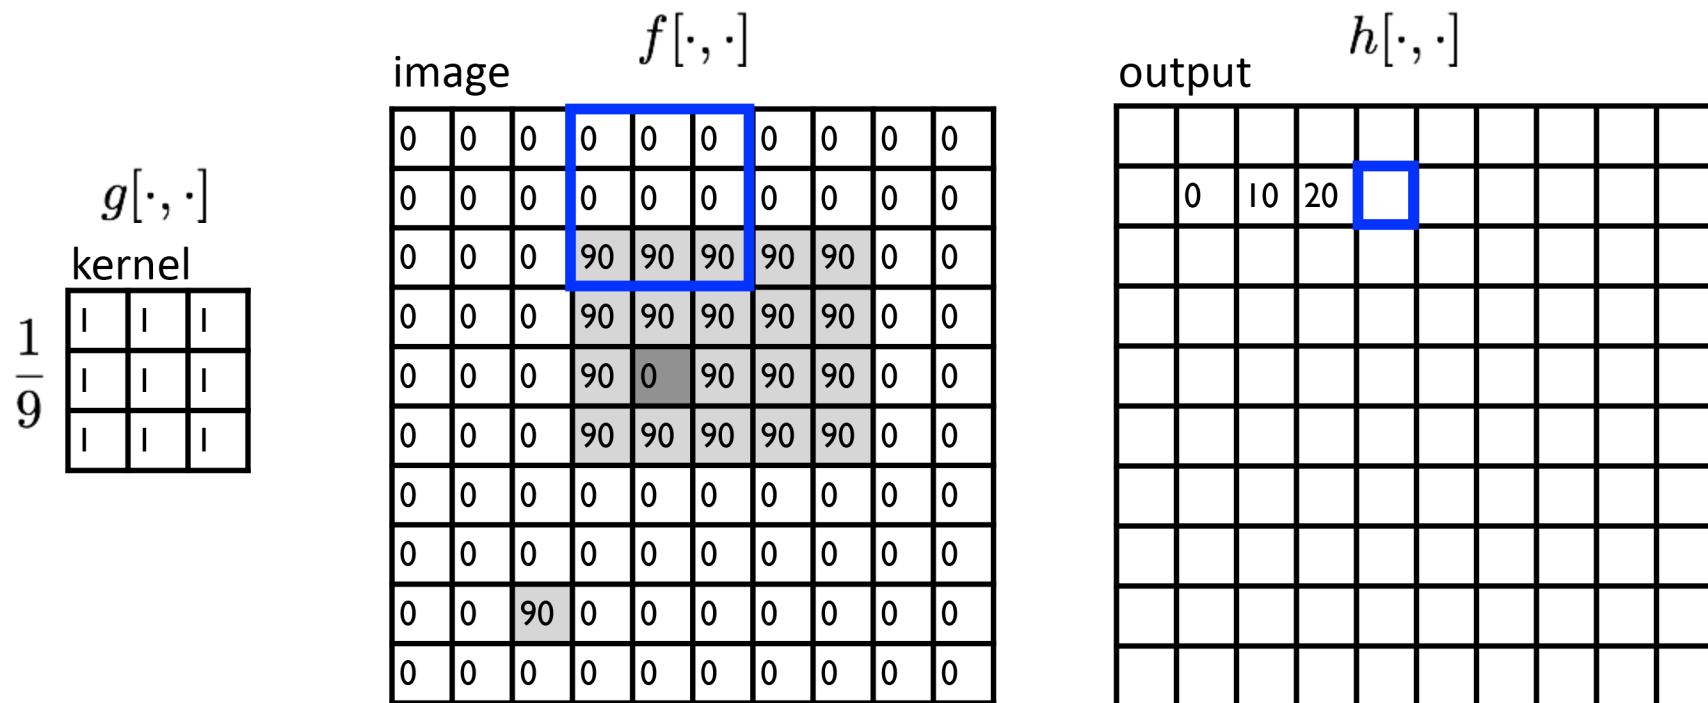
Let's run the box filter



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output filter image (signal)

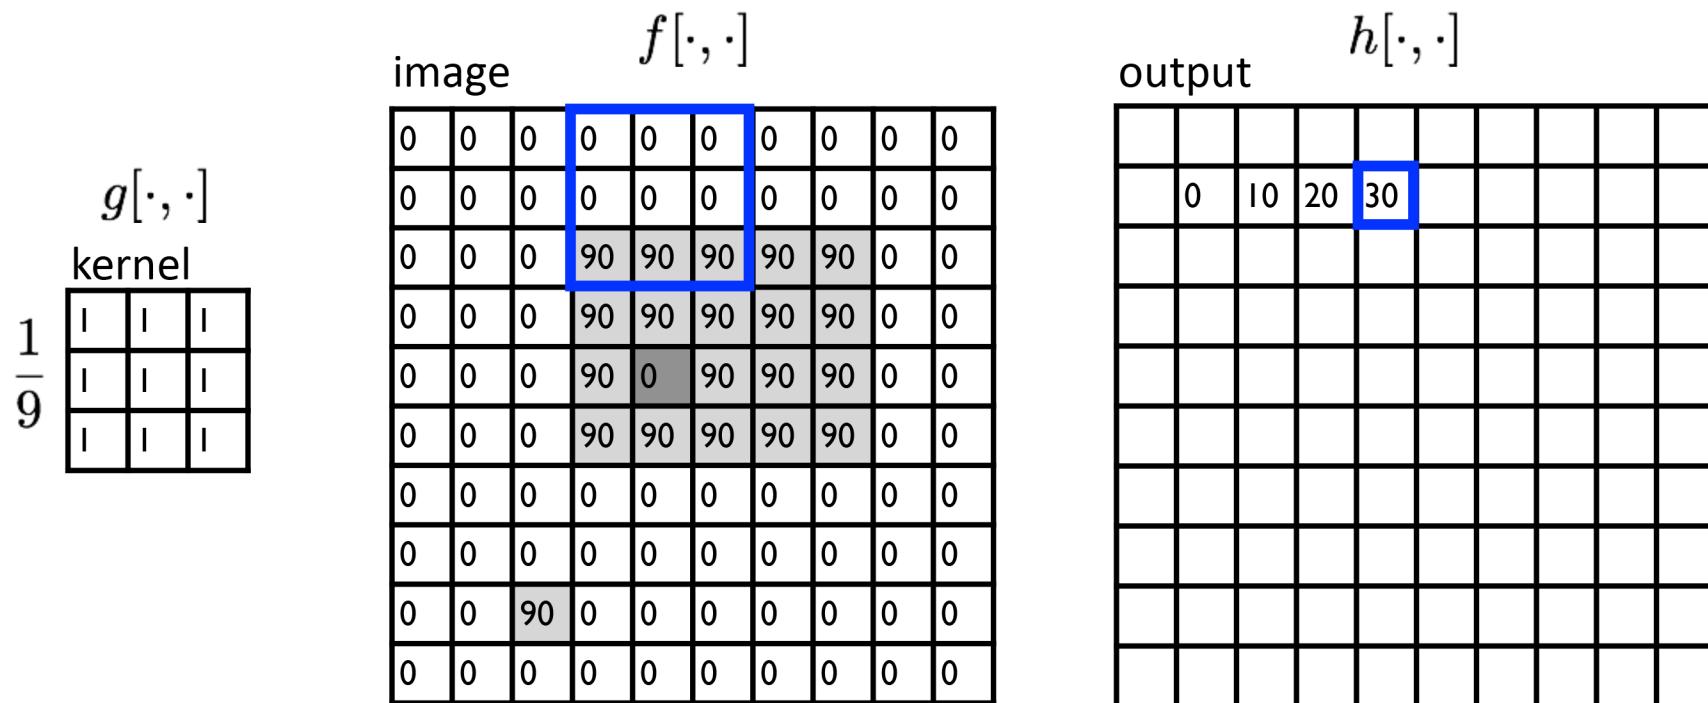
Let's run the box filter



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output filter image (signal)

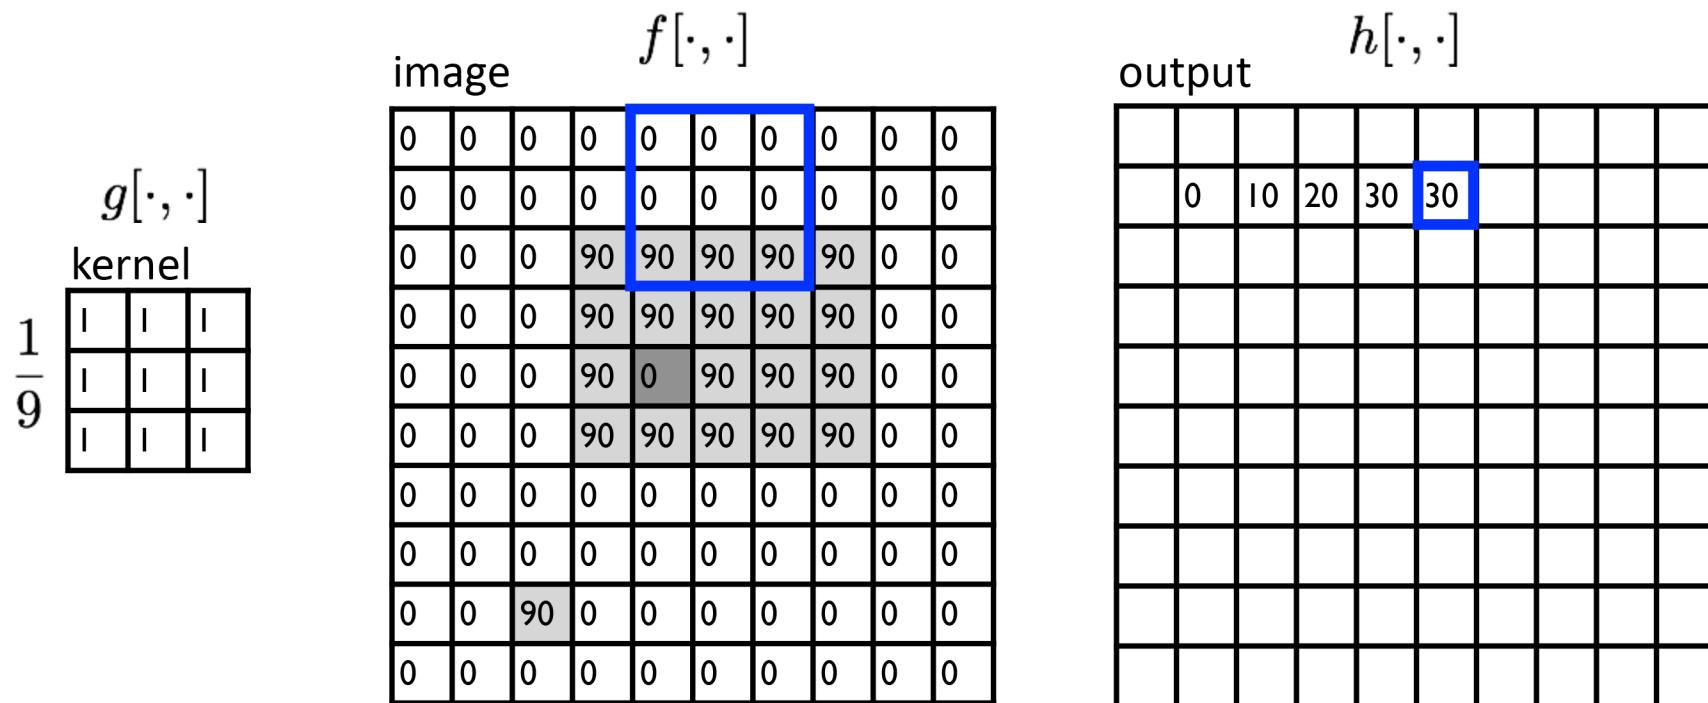
Let's run the box filter



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output filter image (signal)

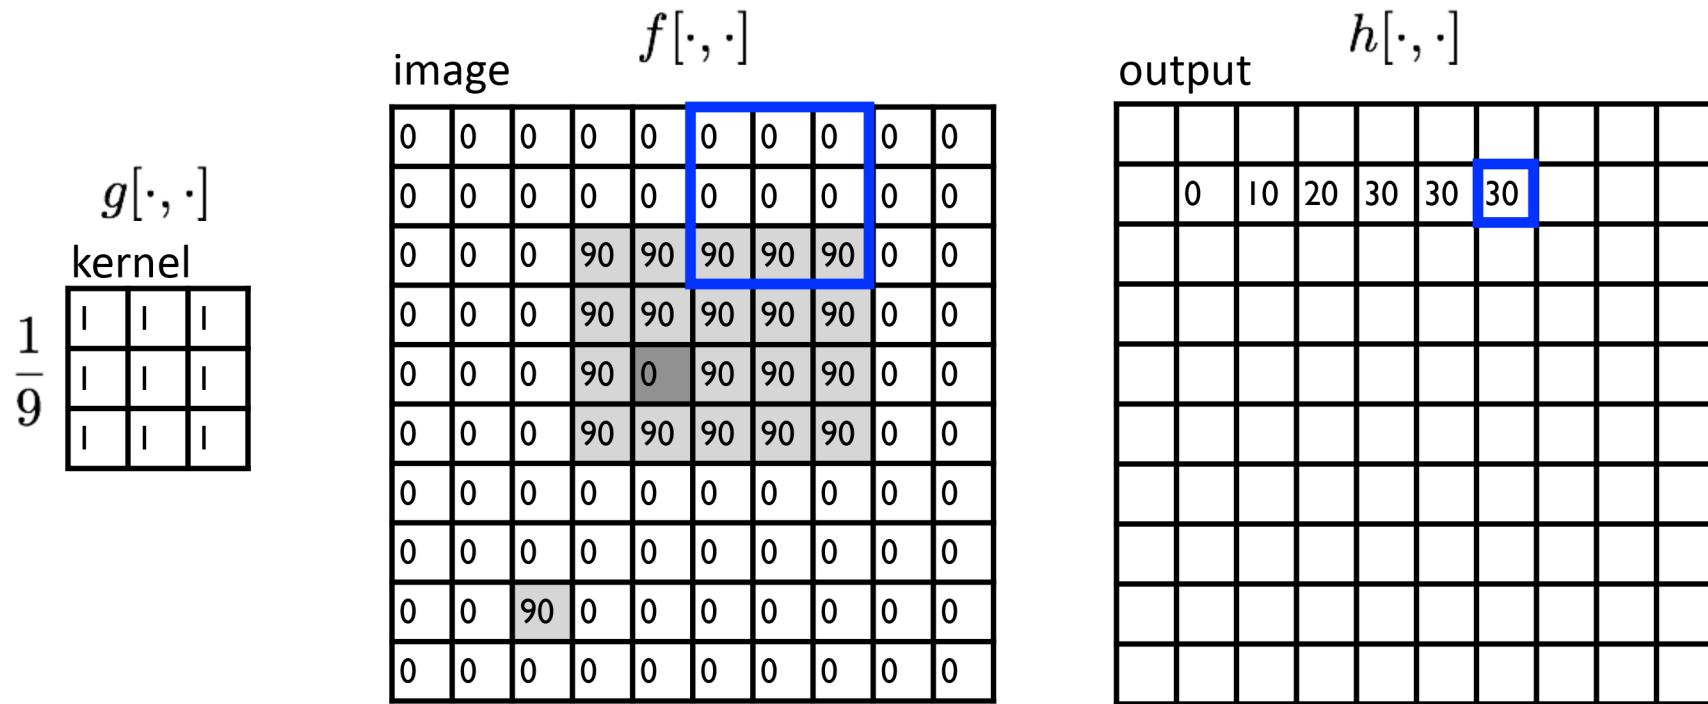
Let's run the box filter



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output filter image (signal)

Let's run the box filter



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output filter image (signal)

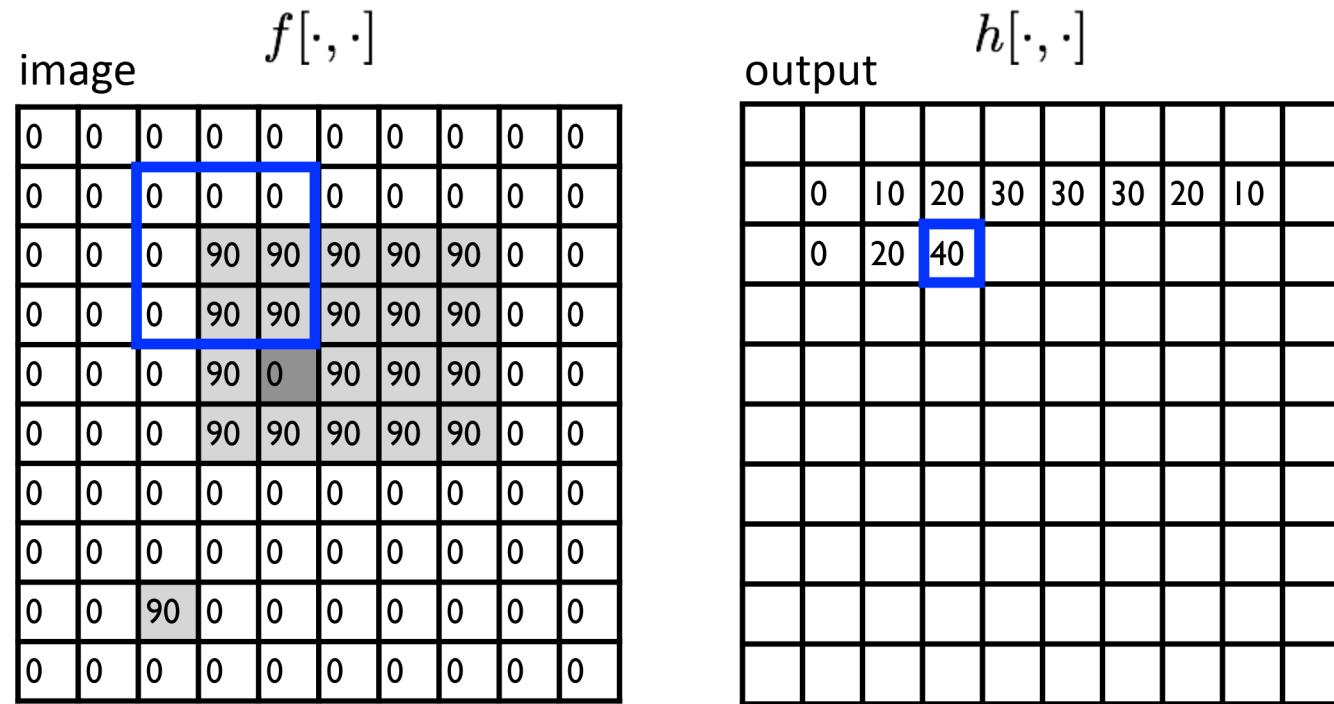
Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9}$$

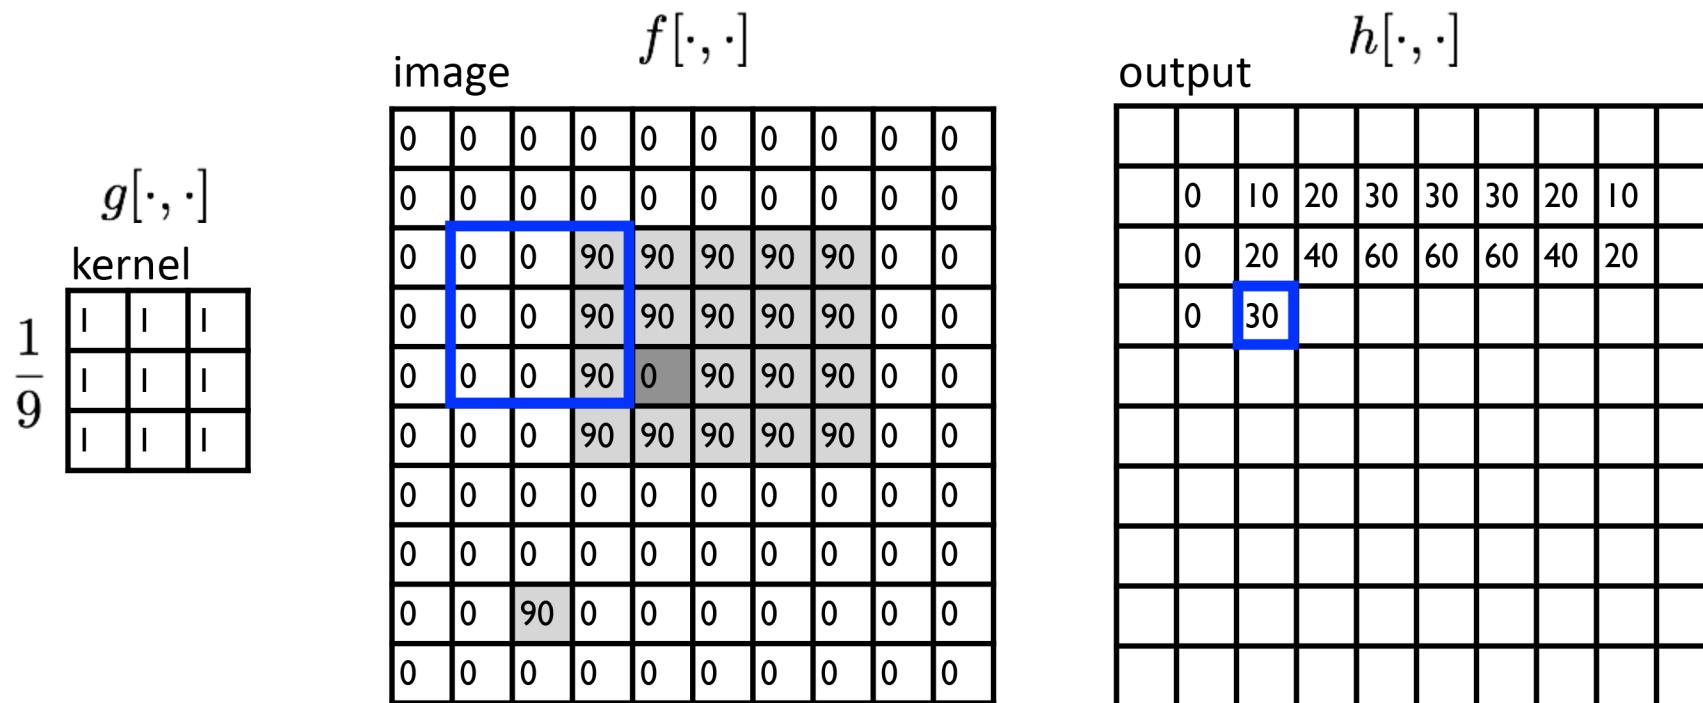
1	1	1
1	1	1
1	1	1



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output filter image (signal)

Let's run the box filter



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output filter image (signal)

Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

image	$f[\cdot, \cdot]$	output	$h[\cdot, \cdot]$
0 0 0 0 0 0 0 0 0 0	0 0 0 0 0 0 0 0 0 0	0 10 20 30 30 30 20 10	0 20 40 60 60 60 40 20
0 0 0 90 90 90 90 90 0 0	0 0 0 90 90 90 90 90 0 0	0 30 50 80 80 90 60 30	0 30 50 80 80 90 60 30
0 0 0 90 90 90 90 90 0 0	0 0 0 90 0 90 90 90 0 0	0 20 30 50 50 60 40 20	0 20 30 50 50 60 40 20
0 0 0 90 0 90 90 90 90 0 0	0 0 0 90 90 90 90 90 0 0	0 10 20 30 30 30 20 10	0 10 20 30 30 30 20 10
0 0 0 0 0 0 0 0 0 0	0 0 0 0 0 0 0 0 0 0	10 10 10 10 0 0 0 0	10 10 10 10 0 0 0 0
0 0 0 0 0 0 0 0 0 0	0 0 0 0 0 0 0 0 0 0	10	

$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output filter image (signal)

Let's run the box filter

$$g[\cdot, \cdot]$$

kernel

$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

image	$f[\cdot, \cdot]$	output	$h[\cdot, \cdot]$
0 0 0 0 0 0 0 0 0 0	0 0 0 0 0 0 0 0 0 0	0 10 20 30 30 30 20 10	0 20 40 60 60 60 40 20
0 0 0 90 90 90 90 90 0 0	0 0 0 90 90 90 90 90 0 0	0 30 50 80 80 90 60 30	0 30 50 80 80 90 60 30
0 0 0 90 90 90 90 90 0 0	0 0 0 90 0 90 90 90 0 0	0 20 30 50 50 60 40 20	0 10 20 30 30 30 20 10
0 0 0 90 0 90 90 90 90 0 0	0 0 0 90 90 90 90 90 0 0	10 10 10 10 0 0 0 0	10 10 10 10 0 0 0 0
0 0 0 0 0 0 0 0 0 0	0 0 0 0 0 0 0 0 0 0	0 0 0 0 0 0 0 0 0 0	0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0	0 0 0 0 0 0 0 0 0 0	0 0 0 0 0 0 0 0 0 0	0 0 0 0 0 0 0 0 0 0
0 0 90 0 0 0 0 0 0 0	0 0 90 0 0 0 0 0 0 0	0 0 0 0 0 0 0 0 0 0	0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0	0 0 0 0 0 0 0 0 0 0	0 0 0 0 0 0 0 0 0 0	0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0	0 0 0 0 0 0 0 0 0 0	0 0 0 0 0 0 0 0 0 0	0 0 0 0 0 0 0 0 0 0

$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output filter image (signal)

... and the result is

Diagram illustrating the convolution process:

Input image $f[\cdot, \cdot]$ (10x10):

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

Kernel $g[\cdot, \cdot]$ (3x3):

$$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

Output $h[\cdot, \cdot]$ (10x10):

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	0	10	20	30	30	30	20	10	
	10	10	10	10	0	0	0	0	
	10	10	10	10	0	0	0	0	

$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

output filter image (signal)

What will be the result of the following convolution operation?



Input image



0	0	0
0	1	0
0	0	0

Filter
(Kernel)

The following convolution operation **produces an identical image**.

Input image $*$ Filter
(Kernel) = Identical image

A diagram illustrating a convolution operation. On the left is a grayscale input image of a person's face. In the center is a multiplication symbol ($*$). To its right is a 3x3 filter kernel with the following values:

0	0	0
0	1	0
0	0	0

Below the filter is the label "Filter (Kernel)". To the right of the filter is an equals sign (=). To the right of the equals sign is another grayscale image of the same person's face, which is identical to the input image. Below this output image is the label "Identical image".

What will be the result of the following convolution operation?



Input image



0	0	0
1	0	0
0	0	0

Filter
(Kernel)

The following convolution operation **shifts the image left by one pixel**.

Input image

*

Filter
(Kernel)

=

Shifted left by 1 pixel

0	0	0
1	0	0
0	0	0

What will be the result of the following convolution operation?



Input image

$$\frac{1}{9}$$

1	1	1
1	1	1
1	1	1

Filter
(Kernel)

The following convolution operation **blurs the image**.

The diagram illustrates a convolution operation. On the left is a grayscale **Input image**, which is a blurry photograph of a person's face. In the center is a **Filter (Kernel)**, represented by a 3x3 grid of orange squares, each containing the value 1. To the left of the filter is a convolution symbol ($*$). To the right of the filter is an equals sign (=). To the right of the equals sign is a grayscale **Blurred image**, which is a much more uniform and less detailed version of the original input image.

$$\text{Input image} \quad * \quad \frac{1}{9} \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array} = \text{Blurred image}$$

Filter
(Kernel)

What will be the result of the following convolution operation?


$$\text{Input image} \quad * \left(\begin{array}{|c|c|c|} \hline 0 & 0 & 0 \\ \hline 0 & 2 & 0 \\ \hline 0 & 0 & 0 \\ \hline \end{array} \right) - \frac{1}{9} \left(\begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array} \right)$$

Filter
(Kernel)

The following convolution operation **sharpens** the image.

The diagram illustrates a convolution operation for image sharpening. It shows the input image, the convolution kernel, the scaling factor, the subtraction term, and the final output image.

Input image:



Kernel:

$$\begin{pmatrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{pmatrix}$$

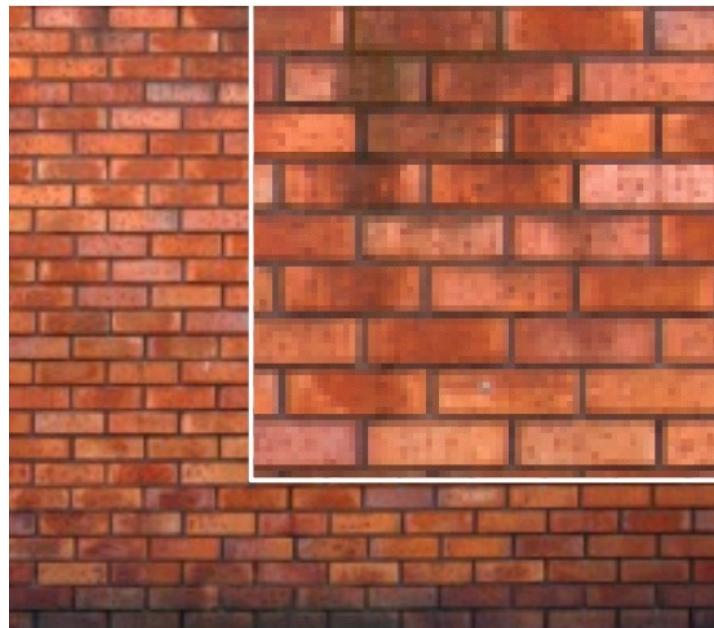
Operations:

- Scale up the intensity
- Remove blurred signals

Final result:

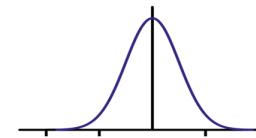

$$\text{Input image} * \left(\begin{pmatrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{pmatrix} - \frac{1}{9} \begin{pmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{pmatrix} \right) = \text{Sharpened image}$$

We can use image filters to [blur images](#), such as using the box filter and the Gaussian filter (2D Gaussian distribution).



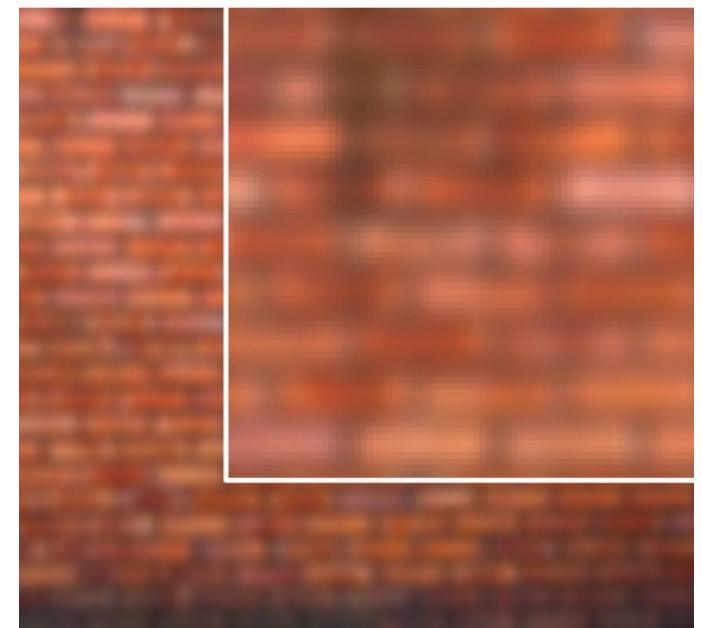
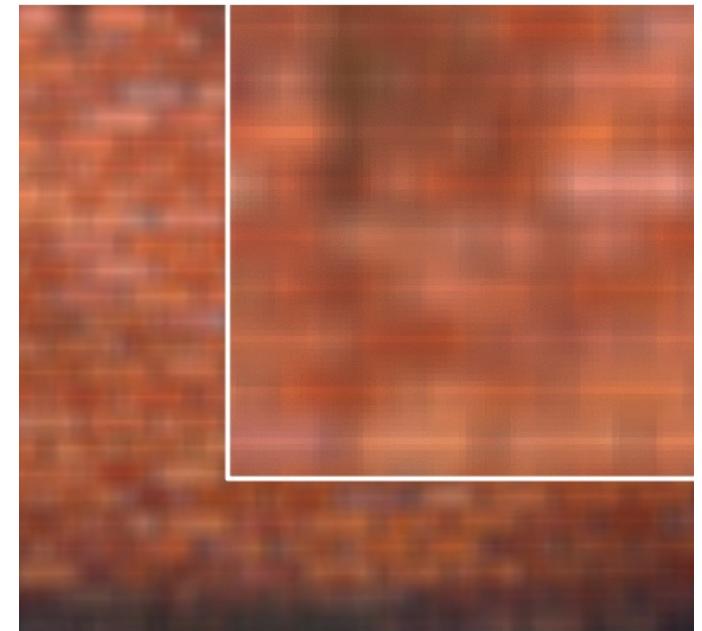
box filter (i.e., mean filter)

$$\frac{1}{9} \times \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

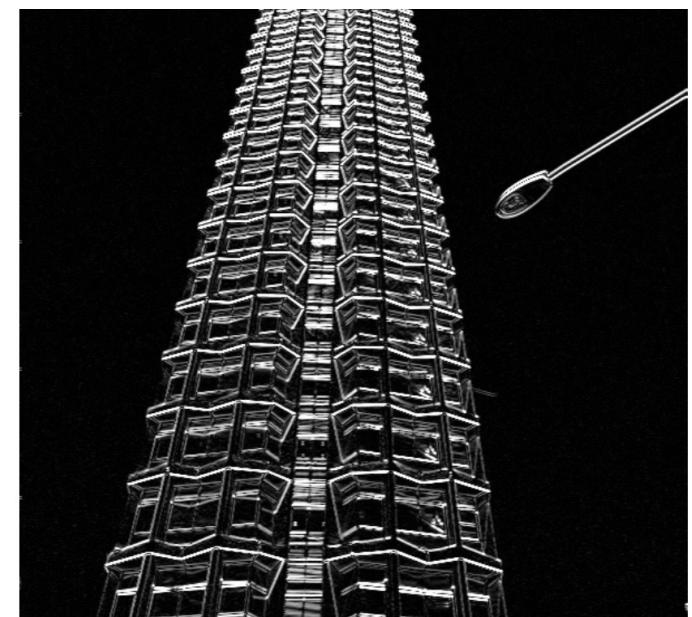
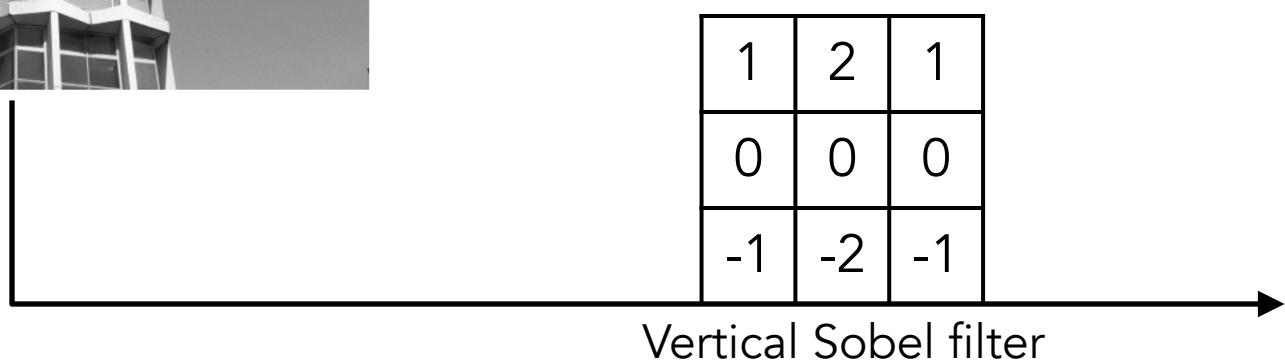
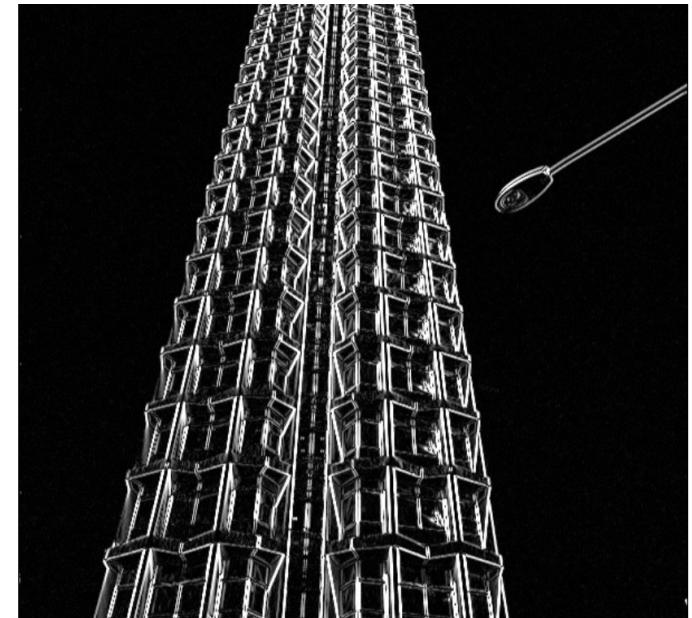
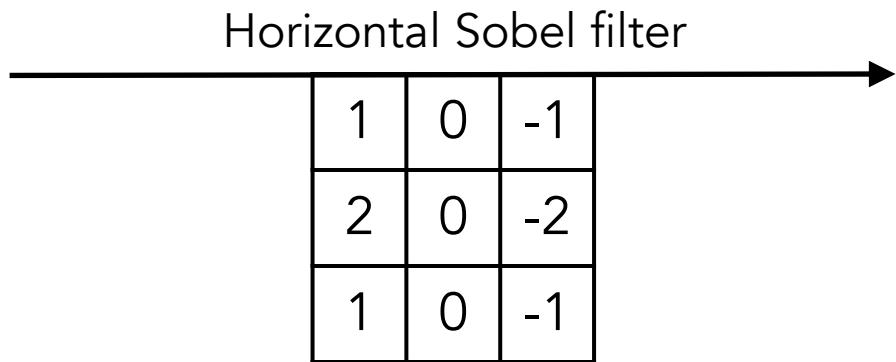


$$\frac{1}{16} \times \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 2 & 4 & 2 \\ \hline 1 & 2 & 1 \\ \hline \end{array}$$

Gaussian filter



We can use image filters to **detect edges**, such as using the Sobel filter to compute the derivative of signals.



We can use various hand-crafted convolutional kernels (i.e., filter banks) to extract different kinds of information in the image, such as using the Gabor filters.



Input image

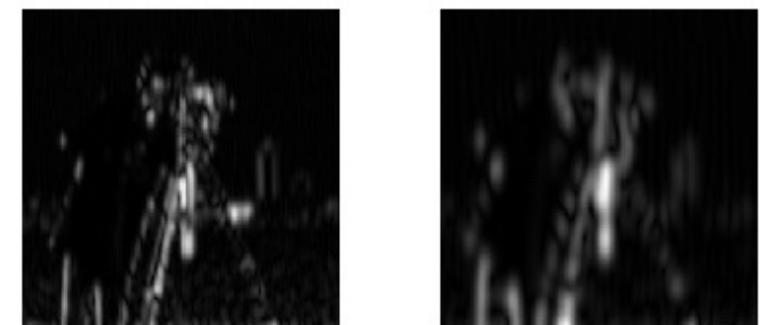


*



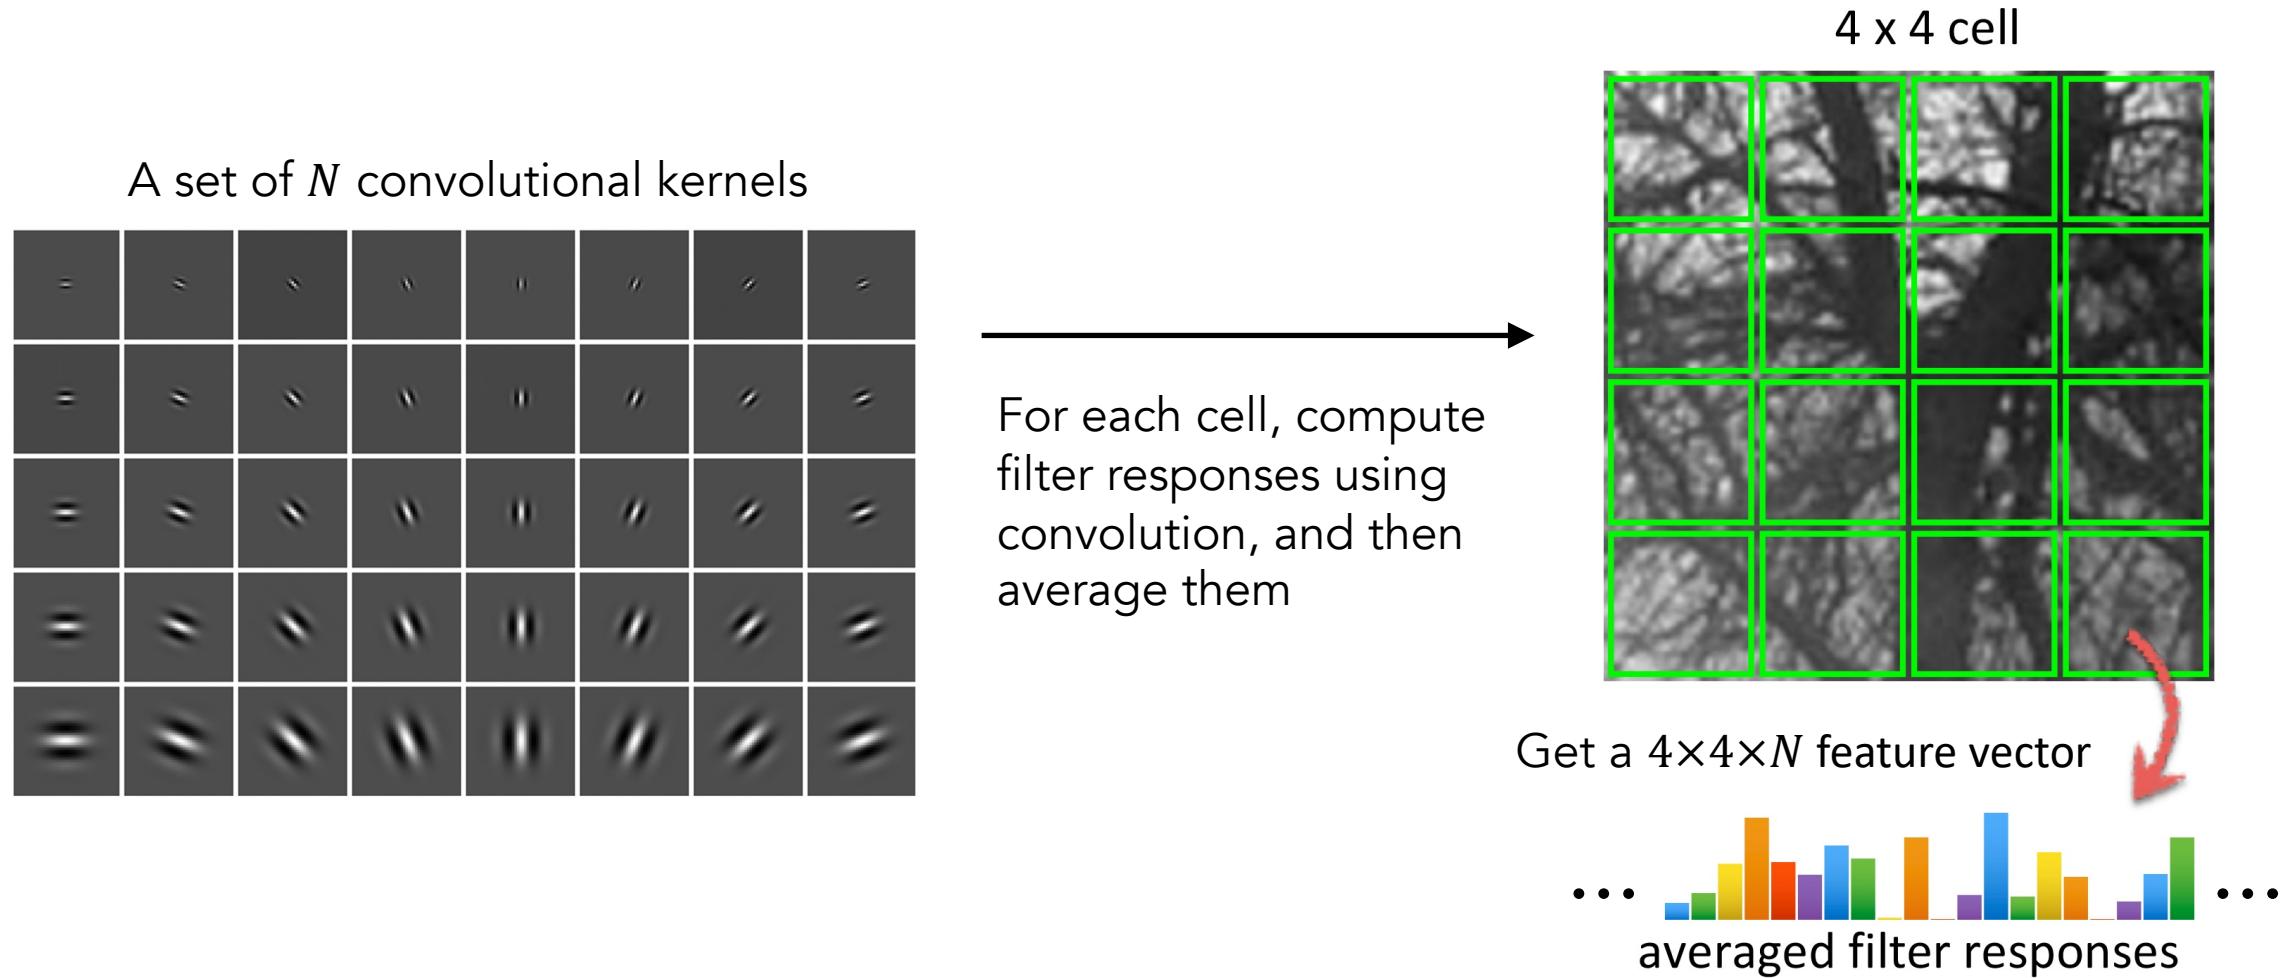
A partial set of Gabor filters

=

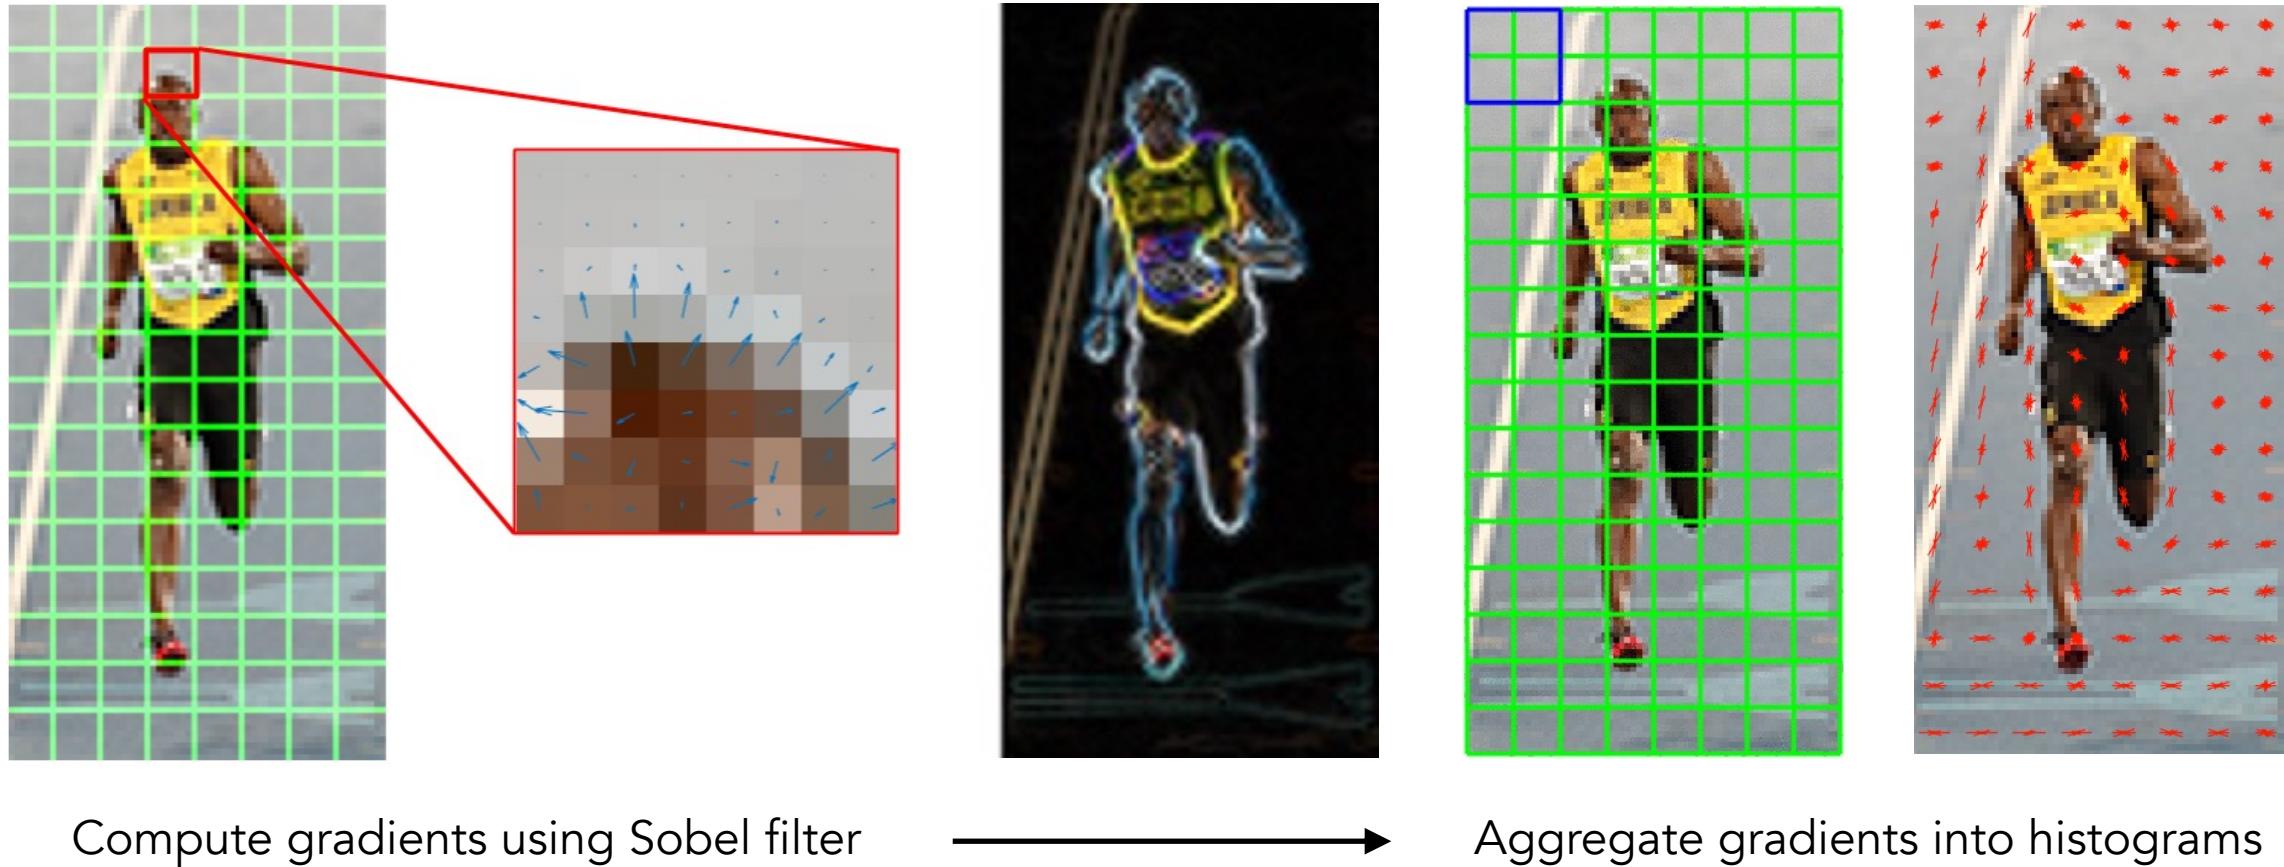


Output images

To construct features, we can perform convolution on image patches using a set of kernels (i.e., filter banks) and then aggregate them into a feature vector.



There are many ways of constructing feature vectors, such as Histogram of Gradients.



Take-Away Messages

- Computers represent images as tensors with numbers in several channels (e.g., red, green, and blue).
- The origin (i.e., index [0,0]) of the image coordinate is at the top-left corner of the image.
- Convolution (the discrete version) means sliding a kernel/filter over the image to compute the sum of element-wise multiplication, which can also be seen as a dot product.
- We can use convolution to perform many different image-filtering operations, such as shifting pixels, blurring/sharpening images, and detecting edges.
- We can use a filter bank (i.e., a collection of kernels) to perform convolution on image patches, and the result (i.e., a set of filter responses) can be used as features.



Questions?